

```
import numpy as np
import pandas as pd
from sklearn.model_selection import train_test_split, cross_val_score
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, classification_report, confusion_matrix
import seaborn as sns
import matplotlib.pyplot as plt

data = pd.read_csv('/content/diabetes_prediction_dataset.csv')
print(data.head())
print(data.info())
print(data.describe())
```



	gender	age	hypertension	heart_disease	smoking_history	bmi	\
0	Female	80.0	0	1	never	25.19	
1	Female	54.0	0	0	No Info	27.32	
2	Male	28.0	0	0	never	27.32	
3	Female	36.0	0	0	current	23.45	
4	Male	76.0	1	1	current	20.14	

	HbA1c_level	blood_glucose_level	diabetes
0	6.6	140	0
1	6.6	80	0
2	5.7	158	0
3	5.0	155	0
4	4.8	155	0

<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 100000 entries, 0 to 99999  
Data columns (total 9 columns):

#	Column	Non-Null Count	Dtype
0	gender	100000 non-null	object
1	age	100000 non-null	float64
2	hypertension	100000 non-null	int64
3	heart_disease	100000 non-null	int64
4	smoking_history	100000 non-null	object
5	bmi	100000 non-null	float64
6	HbA1c_level	100000 non-null	float64
7	blood_glucose_level	100000 non-null	int64
8	diabetes	100000 non-null	int64

dtypes: float64(3), int64(4), object(2)  
memory usage: 6.9+ MB  
None

	age	hypertension	heart_disease	bmi	\
count	100000.000000	100000.000000	100000.000000	100000.000000	
mean	41.885856	0.07485	0.039420	27.320767	
std	22.516840	0.26315	0.194593	6.636783	
min	0.080000	0.00000	0.000000	10.010000	
25%	24.000000	0.00000	0.000000	23.630000	
50%	43.000000	0.00000	0.000000	27.320000	
75%	60.000000	0.00000	0.000000	29.580000	
max	80.000000	1.00000	1.000000	95.690000	

	HbA1c_level	blood_glucose_level	diabetes
count	100000.000000	100000.000000	100000.000000
mean	5.527507	138.058060	0.085000
std	1.070672	40.708136	0.278883
min	3.500000	80.000000	0.000000
25%	4.800000	100.000000	0.000000
50%	5.800000	140.000000	0.000000
75%	6.200000	159.000000	0.000000
max	9.000000	300.000000	1.000000

```

data['blood_glucose_level_category'] = pd.cut(data['blood_glucose_level'], bins=3,

scaler = StandardScaler()
data[['age', 'bmi', 'blood_glucose_level']] = scaler.fit_transform(data[['age', 'bn

X = data[['age', 'bmi', 'hypertension']]
y = data['blood_glucose_level_category']
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_sta

model = LogisticRegression()
model.fit(X_train, y_train)


cv_scores = cross_val_score(model, X_train, y_train, cv=5)
print(f'Cross-validation scores: {cv_scores}')
print(f'Mean cross-validation score: {np.mean(cv_scores)}')

y_pred = model.predict(X_test)
print(f'Accuracy: {accuracy_score(y_test, y_pred)}')
print('Classification Report:')
print(classification_report(y_test, y_pred))
print('Confusion Matrix:')
print(confusion_matrix(y_test, y_pred))

coefficients = pd.DataFrame(model.coef_[0], X.columns, columns=['Coefficient'])
print(coefficients)

sns.heatmap(confusion_matrix(y_test, y_pred), annot=True, fmt='d')
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.show()

```

 Cross-validation scores: [0.58978571 0.5895 0.58964286 0.58985714 0.58985714]  
 Mean cross-validation score: 0.5897285714285714  
 Accuracy: 0.5915666666666667  
 Classification Report:

	precision	recall	f1-score	support
0	0.59	1.00	0.74	17754
1	0.00	0.00	0.00	11470
2	0.05	0.00	0.00	776
accuracy			0.59	30000
macro avg	0.21	0.33	0.25	30000
weighted avg	0.35	0.59	0.44	30000

```
weighted avg      0.00      0.00      0.00      0.00
```

Confusion Matrix:

```
[[17746    0     8]
 [11459    0    11]
 [   775    0     1]]
```

Coefficient

age -0.354876

bmi -0.183636

hypertension -0.259470

```
/usr/local/lib/python3.10/dist-packages/sklearn/metrics/_classification.py:134
```

```
_warn_prf(average, modifier, msg_start, len(result))
```

```
/usr/local/lib/python3.10/dist-packages/sklearn/metrics/_classification.py:134
```

```
_warn_prf(average, modifier, msg_start, len(result))
```

```
/usr/local/lib/python3.10/dist-packages/sklearn/metrics/_classification.py:134
```

```
_warn_prf(average, modifier, msg_start, len(result))
```

