

```
!pip install pandas scikit-learn matplotlib
!pip install pandas matplotlib seaborn
```

```
⇒ Requirement already satisfied: pandas in /usr/local/lib/python3.10/dist-packag
Requirement already satisfied: scikit-learn in /usr/local/lib/python3.10/dist-
Requirement already satisfied: matplotlib in /usr/local/lib/python3.10/dist-pa
Requirement already satisfied: python-dateutil>=2.8.2 in /usr/local/lib/python
Requirement already satisfied: pytz>=2020.1 in /usr/local/lib/python3.10/dist-
Requirement already satisfied: tzdata>=2022.1 in /usr/local/lib/python3.10/di
Requirement already satisfied: numpy>=1.21.0 in /usr/local/lib/python3.10/di
Requirement already satisfied: scipy>=1.3.2 in /usr/local/lib/python3.10/di
Requirement already satisfied: joblib>=1.1.1 in /usr/local/lib/python3.10/di
Requirement already satisfied: threadpoolctl>=2.0.0 in /usr/local/lib/python3
Requirement already satisfied: contourpy>=1.0.1 in /usr/local/lib/python3.10/c
Requirement already satisfied: cycler>=0.10 in /usr/local/lib/python3.10/di
Requirement already satisfied: fonttools>=4.22.0 in /usr/local/lib/python3.10,
Requirement already satisfied: kiwisolver>=1.0.1 in /usr/local/lib/python3.10,
Requirement already satisfied: packaging>=20.0 in /usr/local/lib/python3.10/d
Requirement already satisfied: pillow>=6.2.0 in /usr/local/lib/python3.10/di
Requirement already satisfied: pyparsing>=2.3.1 in /usr/local/lib/python3.10/c
Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.10/dist-pacl
Requirement already satisfied: pandas in /usr/local/lib/python3.10/dist-packag
Requirement already satisfied: matplotlib in /usr/local/lib/python3.10/dist-pa
Requirement already satisfied: seaborn in /usr/local/lib/python3.10/dist-packa
Requirement already satisfied: python-dateutil>=2.8.2 in /usr/local/lib/python
Requirement already satisfied: pytz>=2020.1 in /usr/local/lib/python3.10/dist-
Requirement already satisfied: tzdata>=2022.1 in /usr/local/lib/python3.10/di
Requirement already satisfied: numpy>=1.21.0 in /usr/local/lib/python3.10/di
Requirement already satisfied: contourpy>=1.0.1 in /usr/local/lib/python3.10/c
Requirement already satisfied: cycler>=0.10 in /usr/local/lib/python3.10/di
Requirement already satisfied: fonttools>=4.22.0 in /usr/local/lib/python3.10,
Requirement already satisfied: kiwisolver>=1.0.1 in /usr/local/lib/python3.10,
Requirement already satisfied: packaging>=20.0 in /usr/local/lib/python3.10/d
Requirement already satisfied: pillow>=6.2.0 in /usr/local/lib/python3.10/di
Requirement already satisfied: pyparsing>=2.3.1 in /usr/local/lib/python3.10/c
Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.10/dist-pacl
```

```

import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split, cross_val_score, Stratified
from sklearn.preprocessing import OneHotEncoder
from sklearn.impute import SimpleImputer
from sklearn.tree import DecisionTreeClassifier
from sklearn.model_selection import cross_val_score, cross_val_predict
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_sco
import matplotlib.pyplot as plt
import seaborn as sns

```

```

diabetes_data = pd.read_csv('/content/diabetes_prediction_dataset.csv')
print(diabetes_data.head())

```

```

↔
  gender  age  hypertension  heart_disease  smoking_history  bmi  \
0  Female  80.0             0             1             never  25.19
1  Female  54.0             0             0             No Info  27.32
2   Male   28.0             0             0             never  27.32
3  Female  36.0             0             0             current  23.45
4   Male   76.0             1             1             current  20.14

  HbA1c_level  blood_glucose_level  diabetes
0           6.6                  140         0
1           6.6                   80         0
2           5.7                  158         0
3           5.0                  155         0
4           4.8                  155         0

```

```
# Separate features and target
X = diabetes_data.drop('diabetes', axis=1)
y = diabetes_data['diabetes']

# Identify categorical and numerical columns
categorical_cols = ['gender', 'smoking_history']
numerical_cols = X.columns.difference(categorical_cols)

# Preprocess categorical variables
encoder = OneHotEncoder(sparse=False, handle_unknown='ignore')
encoded_categorical_cols = pd.DataFrame(encoder.fit_transform(X[categorical_cols])
encoded_categorical_cols.columns = encoder.get_feature_names_out(categorical_cols)

# Preprocess numerical variables
imputer = SimpleImputer(strategy='median')
imputed_numerical_cols = pd.DataFrame(imputer.fit_transform(X[numerical_cols]))
imputed_numerical_cols.columns = numerical_cols

# Combine preprocessed categorical and numerical columns
X_processed = pd.concat([encoded_categorical_cols, imputed_numerical_cols], axis=1)

# Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X_processed, y, test_size=0.3)

🔗 /usr/local/lib/python3.10/dist-packages/sklearn/preprocessing/_encoders.py:866:
  warnings.warn(

# Initialize Decision Tree classifier
decision_tree = DecisionTreeClassifier(random_state=42)

# Cross-validation predictions
y_pred_cv = cross_val_predict(decision_tree, X_train, y_train, cv=5)


# Calculate metrics
accuracy = accuracy_score(y_train, y_pred_cv)
precision = precision_score(y_train, y_pred_cv)
recall = recall_score(y_train, y_pred_cv)
f1 = f1_score(y_train, y_pred_cv)
roc_auc = roc_auc_score(y_train, y_pred_cv)

# Fit the model to the full training data and get predictions on the test set
decision_tree.fit(X_train, y_train)
y_pred_test = decision_tree.predict(X_test)
y_pred_proba_test = decision_tree.predict_proba(X_test)[: , 1]
```

```
# ROC curve
fpr, tpr, _ = roc_curve(y_test, y_pred_proba_test)

# Plot ROC curve
plt.figure()
plt.plot(fpr, tpr, label=f'Decision Tree (AUC = {roc_auc:.2f})')
plt.plot([0, 1], [0, 1], 'k--')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC Curve')
plt.legend()
plt.show()

# Display metrics
print("Accuracy:", accuracy)
print("Precision:", precision)
print("Recall:", recall)
print("F1 Score:", f1)
print("ROC AUC:", roc_auc)
```



Accuracy: 0.9515
Precision: 0.7045745361484325
Recall: 0.7399630438434402
F1 Score: 0.721835313396149
ROC AUC: 0.8555624234471624

Suggested code may be subject to a license | 5DcOOKIE/FaceGenius | Shivanshudeveloper/python_url_classification
Start coding or [generate](#) with AI.



	blood_glucose_level	age	bmi	hypertension
blood_glucose_level	1.000000	0.110672	0.091261	0.084429
age	0.110672	1.000000	0.337396	0.251171
bmi	0.091261	0.337396	1.000000	0.147666
hypertension	0.084429	0.251171	0.147666	1.000000

